

```
In [1]: import numpy as np
import pandas as pd
import scipy.stats as stats
import matplotlib.pyplot as plt
from sklearn.datasets import load_boston
import seaborn as sns
from sklearn.model_selection import train_test_split
from statsmodels.stats.outliers_influence import variance_inflation_factor
from sklearn.linear_model import Ridge, Lasso, RidgeCV, LassoCV, ElasticNet, ElasticNetCV, LinearRegression
import statsmodels.formula.api as smf
from sklearn.metrics import mean_squared_error, r2_score
```

```
In [2]: boston = load_boston()
bos = pd.DataFrame(boston.data, columns=boston.feature_names)
```

```
In [3]: bos.head()
```

Out[3]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LST.
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.

```
In [4]: print(boston.keys())

dict_keys(['data', 'target', 'feature_names', 'DESCR', 'filename'])
```

```
In [5]: bos['MEDV'] = boston.target
```

```
In [6]: bos.head()
```

Out[6]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LST.
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.

```
In [7]: bos.isnull().sum()
```

```
Out[7]: CRIM      0  
        ZN        0  
        INDUS    0  
        CHAS     0  
        NOX      0  
        RM       0  
        AGE      0  
        DIS      0  
        RAD      0  
        TAX      0  
        PTRATIO  0  
        B        0  
        LSTAT    0  
        MEDV     0  
        dtype: int64
```

There are no missing values

In [8]: `print(boston.DESCR)`

```
.. _boston_dataset:
```

Boston house prices dataset

-----

**\*\*Data Set Characteristics:\*\***

:Number of Instances: 506

:Number of Attributes: 13 numeric/categorical predictive. Median Value (attribute 14) is usually the target.

:Attribute Information (in order):

- CRIM	per capita crime rate by town
- ZN	proportion of residential land zoned for lots over 25,000 sq.ft.
- INDUS	proportion of non-retail business acres per town
- CHAS	Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)
- NOX	nitric oxides concentration (parts per 10 million)
- RM	average number of rooms per dwelling
- AGE	proportion of owner-occupied units built prior to 1940
- DIS	weighted distances to five Boston employment centres
- RAD	index of accessibility to radial highways
- TAX	full-value property-tax rate per \$10,000
- PTRATIO	pupil-teacher ratio by town
- B	$1000(B_k - 0.63)^2$ where $B_k$ is the proportion of blacks by town
- LSTAT	% lower status of the population
- MEDV	Median value of owner-occupied homes in \$1000's

:Missing Attribute Values: None

:Creator: Harrison, D. and Rubinfeld, D.L.

This is a copy of UCI ML housing dataset.

<https://archive.ics.uci.edu/ml/machine-learning-databases/housing/>

This dataset was taken from the StatLib library which is maintained at Carnegie Mellon University.

The Boston house-price data of Harrison, D. and Rubinfeld, D.L. 'Hedonic prices and the demand for clean air', J. Environ. Economics & Management, vol.5, 81-102, 1978. Used in Belsley, Kuh & Welsch, 'Regression diagnostics ...', Wiley, 1980. N.B. Various transformations are used in the table on pages 244-261 of the latter.

The Boston house-price data has been used in many machine learning papers that address regression problems.

```
.. topic:: References
```

- Belsley, Kuh & Welsch, 'Regression diagnostics: Identifying Influential Data and Sources of Collinearity', Wiley, 1980. 244-261.
- Quinlan, R. (1993). Combining Instance-Based and Model-Based Learning. In

Proceedings on the Tenth International Conference of Machine Learning, 236-243, University of Massachusetts, Amherst. Morgan Kaufmann.

In [9]: `bos.describe()`

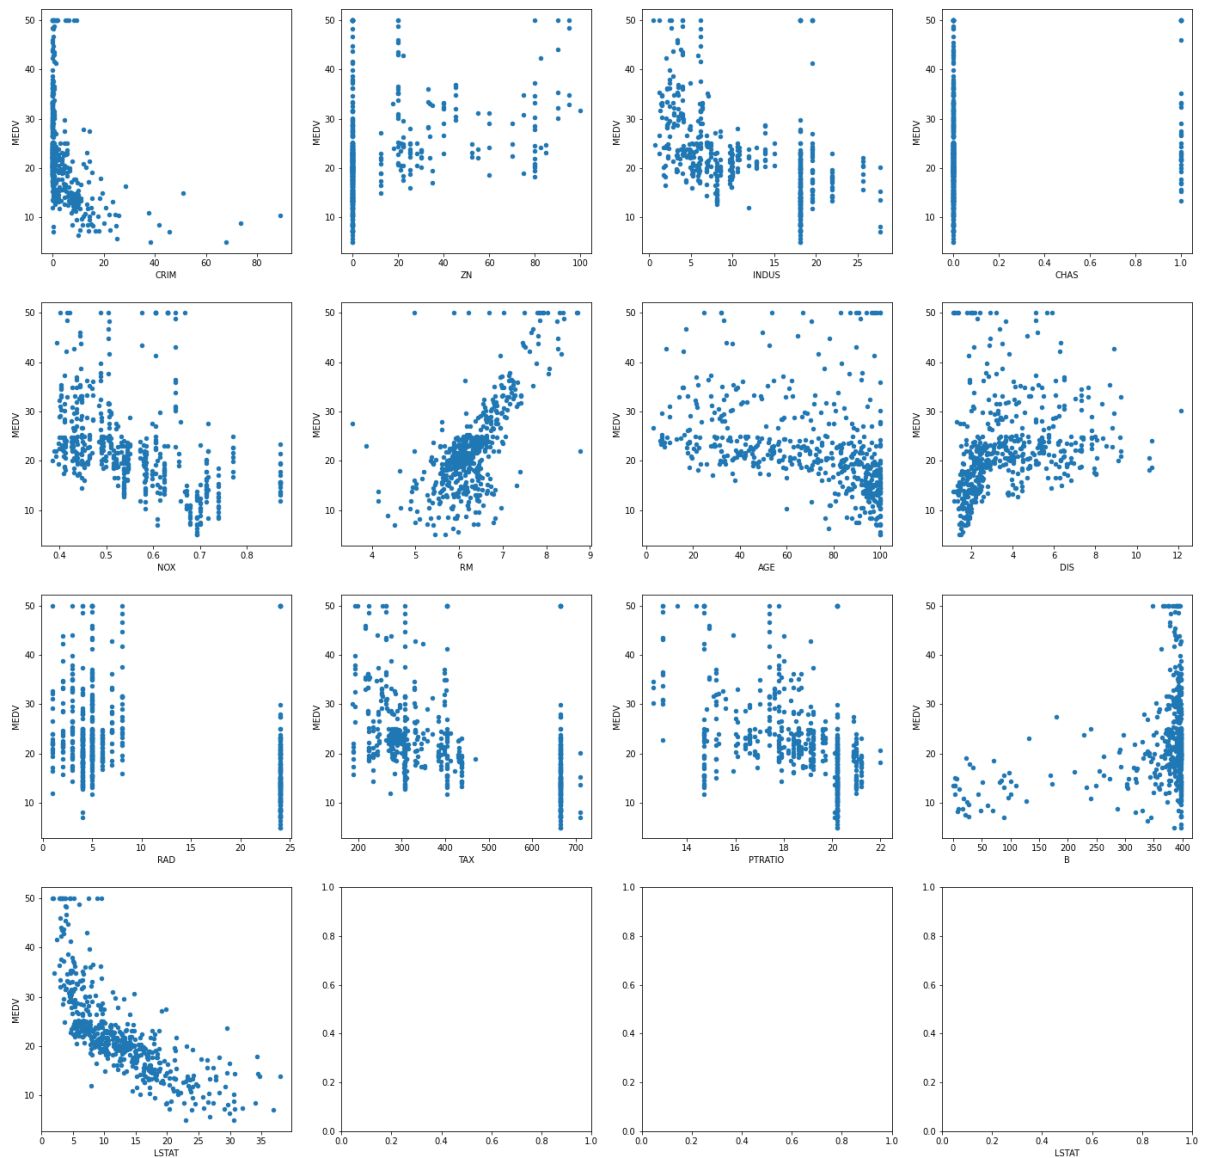
Out[9]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	
<b>count</b>	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506
<b>mean</b>	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634	68.574901	3
<b>std</b>	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617	28.148861	2
<b>min</b>	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	2.900000	1
<b>25%</b>	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500	45.025000	2
<b>50%</b>	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500	77.500000	3
<b>75%</b>	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500	94.075000	5
<b>max</b>	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	100.000000	12

In [10]: `bos.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  ---
0    CRIM        506 non-null    float64
1    ZN          506 non-null    float64
2    INDUS       506 non-null    float64
3    CHAS        506 non-null    float64
4    NOX         506 non-null    float64
5    RM          506 non-null    float64
6    AGE         506 non-null    float64
7    DIS         506 non-null    float64
8    RAD         506 non-null    float64
9    TAX         506 non-null    float64
10   PTRATIO     506 non-null    float64
11   B           506 non-null    float64
12   LSTAT       506 non-null    float64
13   MEDV        506 non-null    float64
dtypes: float64(14)
memory usage: 55.5 KB
```


```
In [11]: # visualize the relationship between the features and the response using scatt
erplots
fig, axs = plt.subplots(4, 4, figsize=(24, 24))
# unpack all the axes subplots
axe = axs.ravel()
for i in range(len(bos.drop(columns='MEDV').columns)):
    bos.plot(kind='scatter', x=bos.columns[i], y='MEDV', ax=ax[i])
    plt.xlabel(bos.columns[i])
```



```
In [12]: sns.set(rc={'figure.figsize':(11.7,8.27)})  
sns.distplot(bos['MEDV'], bins=30)  
plt.show()
```

C:\Users\Urvi\AppData\Roaming\Python\Python37\site-packages\seaborn\distributions.py:2551: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)



```
In [13]: correlation_matrix = bos.corr().round(2)
# annot = True to print the values inside the square
sns.heatmap(data=correlation_matrix, annot=True)
```

Out[13]: <AxesSubplot:>



As seen in correlation matrix:

- LSTAT, PTRATIO, RM are highly correlated with MEDV
- TAX and RAD are highly correlated
- INDUS, NOX and AGE are high correlated with DIS
- INDUS is highly correlated with ZN, TAX, NOX

```
In [14]: X = bos.drop(columns='MEDV')
y = bos.MEDV
```

```
In [15]: X.shape[1]
```

Out[15]: 13

```
In [16]: ## Checking vif
vif = pd.DataFrame()
vif["VIF"] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
vif["Features"] = X.columns
```



In [17]: vif

Out[17]:

	VIF	Features
0	2.100373	CRIM
1	2.844013	ZN
2	14.485758	INDUS
3	1.152952	CHAS
4	73.894947	NOX
5	77.948283	RM
6	21.386850	AGE
7	14.699652	DIS
8	15.167725	RAD
9	61.227274	TAX
10	85.029547	PTRATIO
11	20.104943	B
12	11.102025	LSTAT

Observations in VIF:

- INDUS, NOX, RM, AGE, DIS, RAD, TAX, PTRATIO, B, LSTAT have vif > 10

```
In [18]: ## Dropping TAX, as tax and rad had high correlation
X = bos.drop(columns=['MEDV', 'TAX'])
y = bos.MEDV
## Checking vif
vif = pd.DataFrame()
vif["VIF"] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
vif["Features"] = X.columns
vif
```

Out[18]:

	VIF	Features
0	2.100323	CRIM
1	2.697230	ZN
2	11.743319	INDUS
3	1.136630	CHAS
4	71.972959	NOX
5	77.946536	RM
6	21.377489	AGE
7	14.641579	DIS
8	5.599479	RAD
9	82.355181	PTRATIO
10	20.104332	B
11	11.098281	LSTAT

```
In [19]: ## Dropping NOX, as Nox and Dis, indus had high correlation
X = bos.drop(columns=['MEDV', 'TAX', 'NOX'])
y = bos.MEDV
## Checking vif
vif = pd.DataFrame()
vif["VIF"] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
vif["Features"] = X.columns
vif
```

Out[19]:

	VIF	Features
0	2.098376	CRIM
1	2.696984	ZN
2	9.930116	INDUS
3	1.134868	CHAS
4	58.059674	RM
5	19.826168	AGE
6	14.485117	DIS
7	5.405429	RAD
8	82.299026	PTRATIO
9	19.872129	B
10	10.116939	LSTAT

```
In [20]: ## Dropping other columns with high correlation as well
X = bos.drop(columns=['MEDV', 'TAX', 'NOX', 'B', 'AGE', 'DIS'])
y = bos.MEDV
## Checking vif
vif = pd.DataFrame()
vif["VIF"] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
vif["Features"] = X.columns
vif
```

Out[20]:

	VIF	Features
0	2.048566	CRIM
1	1.801368	ZN
2	8.164547	INDUS
3	1.123261	CHAS
4	43.767132	RM
5	4.867376	RAD
6	57.716508	PTRATIO
7	8.061584	LSTAT

```
In [21]: x_train,x_test,y_train,y_test = train_test_split(X, y, test_size = 0.25, random_state=355)
```

```
In [22]: regression = LinearRegression()
         regression.fit(x_train, y_train)
```

```
Out[22]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

```
In [23]: regression.score(x_train, y_train)
```

```
Out[23]: 0.6852743389124581
```

```
In [24]: # Let's create a function to create adjusted R-Squared
         def adj_r2(x,y):
             r2 = regression.score(x,y)
             n = x.shape[0]
             p = x.shape[1]
             adjusted_r2 = 1-(1-r2)*(n-1)/(n-p-1)
             return adjusted_r2
```

```
In [25]: adj_r2(x_train,y_train)
```

```
Out[25]: 0.6784694597538086
```

```
In [26]: lm = smf.ols(formula='MEDV ~ CRIM + ZN + INDUS + CHAS + RM + RAD + PTRATIO + LSTAT', data=bos).fit()
         lm.conf_int()
```

```
Out[26]:
```

	0	1
<b>Intercept</b>	10.619480	27.102404
<b>CRIM</b>	-0.157797	-0.019711
<b>ZN</b>	-0.029286	0.018326
<b>INDUS</b>	-0.137108	0.061141
<b>CHAS</b>	1.523579	5.150357
<b>RM</b>	3.549548	5.242013
<b>RAD</b>	-0.012834	0.147742
<b>PTRATIO</b>	-1.179944	-0.668984
<b>LSTAT</b>	-0.647857	-0.452509

In [27]: lm.summary()

Out[27]: OLS Regression Results

<b>Dep. Variable:</b>	MEDV	<b>R-squared:</b>	0.692
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.687
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	139.5
<b>Date:</b>	Fri, 15 Jan 2021	<b>Prob (F-statistic):</b>	7.94e-122
<b>Time:</b>	22:25:48	<b>Log-Likelihood:</b>	-1542.4
<b>No. Observations:</b>	506	<b>AIC:</b>	3103.
<b>Df Residuals:</b>	497	<b>BIC:</b>	3141.
<b>Df Model:</b>	8		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>Intercept</b>	18.8609	4.195	4.496	0.000	10.619	27.102
<b>CRIM</b>	-0.0888	0.035	-2.526	0.012	-0.158	-0.020
<b>ZN</b>	-0.0055	0.012	-0.452	0.651	-0.029	0.018
<b>INDUS</b>	-0.0380	0.050	-0.753	0.452	-0.137	0.061
<b>CHAS</b>	3.3370	0.923	3.615	0.000	1.524	5.150
<b>RM</b>	4.3958	0.431	10.206	0.000	3.550	5.242
<b>RAD</b>	0.0675	0.041	1.651	0.099	-0.013	0.148
<b>PTRATIO</b>	-0.9245	0.130	-7.110	0.000	-1.180	-0.669
<b>LSTAT</b>	-0.5502	0.050	-11.067	0.000	-0.648	-0.453

<b>Omnibus:</b>	184.679	<b>Durbin-Watson:</b>	0.978
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	851.037
<b>Skew:</b>	1.567	<b>Prob(JB):</b>	1.58e-185
<b>Kurtosis:</b>	8.527	<b>Cond. No.</b>	579.

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

As p values for ZN and INDUS > 0.05, we fail to reject null hypothesis. Hence, there is no correlation between MEDV and ZN, INDUS. So, we drop them

```
In [28]: lm = smf.ols(formula='MEDV ~ CRIM + CHAS + RM + RAD + PTRATIO + LSTAT', data=boston).fit()
lm.summary()
```

Out[28]: OLS Regression Results

<b>Dep. Variable:</b>	MEDV	<b>R-squared:</b>	0.691
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.688
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	186.4
<b>Date:</b>	Fri, 15 Jan 2021	<b>Prob (F-statistic):</b>	5.75e-124
<b>Time:</b>	22:25:48	<b>Log-Likelihood:</b>	-1542.7
<b>No. Observations:</b>	506	<b>AIC:</b>	3099.
<b>Df Residuals:</b>	499	<b>BIC:</b>	3129.
<b>Df Model:</b>	6		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>Intercept</b>	18.2899	4.086	4.476	0.000	10.262	26.318
<b>CRIM</b>	-0.0885	0.035	-2.527	0.012	-0.157	-0.020
<b>CHAS</b>	3.2827	0.912	3.601	0.000	1.492	5.074
<b>RM</b>	4.4174	0.428	10.316	0.000	3.576	5.259
<b>RAD</b>	0.0567	0.038	1.497	0.135	-0.018	0.131
<b>PTRATIO</b>	-0.9159	0.126	-7.274	0.000	-1.163	-0.668
<b>LSTAT</b>	-0.5583	0.047	-11.950	0.000	-0.650	-0.467

<b>Omnibus:</b>	182.840	<b>Durbin-Watson:</b>	0.979
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	833.465
<b>Skew:</b>	1.553	<b>Prob(JB):</b>	1.04e-181
<b>Kurtosis:</b>	8.467	<b>Cond. No.</b>	478.

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

As p values for RAD > 0.05, we fail to reject null hypothesis. Hence, there is no correlation between MEDV and RAD. So, we drop them

```
In [29]: lm = smf.ols(formula='MEDV ~ CRIM + CHAS + RM + PTRATIO + LSTAT', data=bos).fit()
lm.summary()
```

Out[29]: OLS Regression Results

<b>Dep. Variable:</b>	MEDV	<b>R-squared:</b>	0.690
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.687
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	222.7
<b>Date:</b>	Fri, 15 Jan 2021	<b>Prob (F-statistic):</b>	1.11e-124
<b>Time:</b>	22:25:48	<b>Log-Likelihood:</b>	-1543.9
<b>No. Observations:</b>	506	<b>AIC:</b>	3100.
<b>Df Residuals:</b>	500	<b>BIC:</b>	3125.
<b>Df Model:</b>	5		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>Intercept</b>	16.5736	3.927	4.221	0.000	8.859	24.288
<b>CRIM</b>	-0.0625	0.030	-2.052	0.041	-0.122	-0.003
<b>CHAS</b>	3.3879	0.910	3.723	0.000	1.600	5.176
<b>RM</b>	4.5262	0.423	10.713	0.000	3.696	5.356
<b>PTRATIO</b>	-0.8489	0.118	-7.204	0.000	-1.080	-0.617
<b>LSTAT</b>	-0.5397	0.045	-11.970	0.000	-0.628	-0.451

<b>Omnibus:</b>	198.984	<b>Durbin-Watson:</b>	0.968
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	1006.250
<b>Skew:</b>	1.670	<b>Prob(JB):</b>	3.13e-219
<b>Kurtosis:</b>	9.048	<b>Cond. No.</b>	417.

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [30]: bos.columns
```

```
Out[30]: Index(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD', 'TAX',
              'PTRATIO', 'B', 'LSTAT', 'MEDV'],
              dtype='object')
```

```
In [31]: X = bos.drop(columns=['ZN', 'INDUS', 'NOX', 'AGE', 'DIS', 'RAD', 'TAX', 'B'])
x_train,x_test,y_train,y_test = train_test_split(X, y, test_size = 0.25, random_state=355)
regression = LinearRegression()
regression.fit(x_train, y_train)
regression.score(x_test,y_test)
```

Out[31]: 1.0

```
In [32]: adj_r2(x_test,y_test)
```

Out[32]: 1.0

```
In [33]: # Lasso Regularization
# LassoCV will return best alpha and coefficients after performing 10 cross validations
lasscv = LassoCV(alphas = None,cv =10, max_iter = 100000, normalize = True)
lasscv.fit(x_train, y_train)
```

Out[33]: LassoCV(alphas=None, copy\_X=True, cv=10, eps=0.001, fit\_intercept=True, max\_iter=100000, n\_alphas=100, n\_jobs=None, normalize=True, positive=False, precompute='auto', random\_state=None, selection='cyclic', tol=0.0001, verbose=False)

```
In [34]: # best alpha parameter
alpha = lasscv.alpha_
alpha
```

Out[34]: 0.00046719468953559777

```
In [35]: #now that we have best parameter, let's use Lasso regression and see how well our data has fitted before

lasso_reg = Lasso(alpha)
lasso_reg.fit(x_train, y_train)
```

Out[35]: Lasso(alpha=0.00046719468953559777, copy\_X=True, fit\_intercept=True, max\_iter=1000, normalize=False, positive=False, precompute=False, random\_state=None, selection='cyclic', tol=0.0001, warm\_start=False)

```
In [36]: lasso_reg.score(x_test, y_test)
```

Out[36]: 0.9999999999681303



```
In [37]: # Using Ridge regression model
# RidgeCV will return best alpha and coefficients after performing 10 cross va
lidations.
# We will pass an array of random numbers for ridgeCV to select best alpha fro
m them

alphas = np.random.uniform(low=0, high=10, size=(50,))
ridgecv = RidgeCV(alphas = alphas, cv=10, normalize = True)
ridgecv.fit(x_train, y_train)
```

```
Out[37]: RidgeCV(alphas=array([5.10651279, 0.54103942, 3.81847241, 5.74994119, 8.82230
368,
      3.41799351, 0.16279025, 2.70816071, 5.02751315, 7.62855495,
      9.73078147, 5.13815897, 3.41791352, 8.17535891, 0.23303335,
      9.91519295, 5.2274401 , 3.44794737, 7.71754308, 5.29620168,
      5.47707288, 9.02099261, 1.08240683, 4.90941542, 3.98616556,
      8.25057402, 4.84580458, 4.33898855, 7.11215745, 9.33756954,
      7.5634796 , 6.21268459, 9.79367692, 4.28561885, 1.78504003,
      5.75871426, 5.16825768, 3.27992815, 1.53237995, 2.42598572,
      5.60075164, 1.06301355, 1.66238044, 7.93355354, 6.59636197,
      7.69310213, 8.55628208, 9.37934931, 1.29629281, 9.99743755]),
      cv=10, fit_intercept=True, gcv_mode=None, normalize=True, scoring=Non
e,
      store_cv_values=False)
```

```
In [38]: ridgecv.alpha_
```

```
Out[38]: 0.16279024933434139
```

```
In [39]: ridge_model = Ridge(alpha=ridgecv.alpha_)
ridge_model.fit(x_train, y_train)
```

```
Out[39]: Ridge(alpha=0.16279024933434139, copy_X=True, fit_intercept=True, max_iter=No
ne,
      normalize=False, random_state=None, solver='auto', tol=0.001)
```

```
In [40]: ridge_model.score(x_test, y_test)
```

```
Out[40]: 0.999999999920012
```

As Lasso, Ridge and Linear Regression Test score is almost same, model is not overfitted.

```
In [41]: # Elastic net

elasticCV = ElasticNetCV(alphas = None, cv =10)
elasticCV.fit(x_train, y_train)
```

```
Out[41]: ElasticNetCV(alphas=None, copy_X=True, cv=10, eps=0.001, fit_intercept=True,
      l1_ratio=0.5, max_iter=1000, n_alphas=100, n_jobs=None,
      normalize=False, positive=False, precompute='auto',
      random_state=None, selection='cyclic', tol=0.0001, verbose=0)
```

```
In [42]: elasticCV.alpha_
```

```
Out[42]: 0.1654493254711398
```

```
In [43]: elasticCV.l1_ratio
```

```
Out[43]: 0.5
```

```
In [44]: elasticnet_reg = ElasticNet(alpha=elasticCV.alpha_, l1_ratio=0.5)
elasticnet_reg.fit(x_train, y_train)
```

```
Out[44]: ElasticNet(alpha=0.1654493254711398, copy_X=True, fit_intercept=True,
                    l1_ratio=0.5, max_iter=1000, normalize=False, positive=False,
                    precompute=False, random_state=None, selection='cyclic', tol=0.000
                    1,
                    warm_start=False)
```

```
In [45]: elasticnet_reg.score(x_test, y_test)
```

```
Out[45]: 0.9999959500003314
```

```
In [46]: # model evaluation for training set
y_train_predict = regression.predict(x_train)
rmse = (np.sqrt(mean_squared_error(y_train, y_train_predict)))
r2 = r2_score(y_train, y_train_predict)

print("The model performance for training set")
print("-----")
print('RMSE is {}'.format(rmse))
print('R2 score is {}'.format(r2))
print("\n")

# model evaluation for testing set
y_test_predict = regression.predict(x_test)
rmse = (np.sqrt(mean_squared_error(y_test, y_test_predict)))
r2 = r2_score(y_test, y_test_predict)

print("The model performance for testing set")
print("-----")
print('RMSE is {}'.format(rmse))
print('R2 score is {}'.format(r2))
```

The model performance for training set

-----

RMSE is 5.620667389044694e-15

R2 score is 1.0


The model performance for testing set

-----

RMSE is 6.3428626160647376e-15

R2 score is 1.0

```
In [47]: predicted = regression.predict(x_test)
residuals = y_test-predicted
sns.set(rc={'figure.figsize':(11,8)})
sns.displot(residuals, bins=10)
plt.show()
```



```
In [48]: # mean of residuals is almost 0
residuals.mean()
```

```
Out[48]: -6.224242468764657e-16
```

```
In [49]: ## Hence all the assumptions are satisfied
```

Project Done By: Urvi Gadda

mailto: urvigada96@gmail.com

```
In [ ]:
```